HCLTech
Supercharging Progress™

Analyst Data CoPilot

Agentic AI Use Case Draft Proposal

HCLTech

Contents

1. Analyst Data Copilot (MVP)

1   Overview of the Use Case

This proposal outlines a pragmatic, deliverable-focused MVP which we'll complete as a co-innovation project between **HCLTech** and **ING**. The solution is a **Analyst Data Copilot** designed to accelerate the initial investigation of Financial Crime and Fraud Prevention (FCFP) alerts.

- **Problem**: FCFP L1 analysts spend significant manual effort gathering and reviewing data from multiple systems, primarily FCRM, just to perform an initial triage. This repetitive, time-consuming process slows down case resolution.

- **Solution** (MVP): We'll build an Agentic AI solution. When an analyst provides a Case ID, the Agent's goal is to autonomously gather and summarize all relevant case data.
    1. **Autonomous Data Gathering**: The Agent will use a "tool", which accessed a MS SQL Database, to fetch a predefined set of critical data points about the transaction, customer and accounts from the FCRM system.
    2. **GenAI-Powered Analysis**: The Agent will use a GenAI model (Gemini Pro) to analyze this data, present a concise natural-language summary of the gathered data. This output will not include any decisions, proposals, or recommendations.
- **Value**: This focused MVP provides immediate value by reducing manual work and speeding up decisions. It's designed as a production-quality application, enabling us to move quickly and establish a secure foundation for Agentic AI. The goal is to deliver a solution ready for formal User Acceptance Testing (UAT) at the end of this engagement.

2   Scope Description

To ensure deliverability within the agreed co-innovation budget, the scope is tightly defined.

2.1   MVP (target for UAT)

- **Functionality**:
    - A standalone web application where an analyst can input a Case ID
    - The Agent autonomously triggers a read-only data fetch from the FCRM system only
    - The user interface (UI) displays a "CoPilot" panel with:
        1. A structured view of the gathered data
        2. A natural-language summary of the key facts
- **Out of Scope for MVP**:
    - Any decision-making, recommendations, or proposals (e.g., "Close" or "Escalate")
    - Direct "write-back" integration to FCRM (e.g., no auto-closing cases)
    - Data integration from other sources (e.g., Quantexa, TMA, SharePoint)
    - Fully interactive, natural-language Q&A
    - Automated aggregation or creation of new alerts

2.2   High-level release cycle (post-MVP)

- **Release 2** (Embedded integration): We'll embed the Copilot as a native widget within the FCRM front-end UI. This allows analysts to invoke the Agent's analysis without leaving their primary workflow
- **Release 3 (Unstructured data integration):** We'll integrate unstructured data sources (e.g., policy documents, SharePoint). This allows the Copilot to answer analyst questions about compliance procedures, not just case data
- **Release 4 (Multi-system data integration):** We'll expand the Agent's "tools" to query structured data from additional systems like **Quantexa** and the **TMA Tool**. This enables the "single pane of glass" holistic analysis
- **Release 5 (Agent-in-the-loop actions):** *Pending business and SATM approval,* we would explore enhancing the Agent with "tools" to perform 'write' actions (e.g., auto-populating notes) or suggest recommendations

---

3   19-Week MVP Timeline

The timeline isn't driven by the complexity of the UI, but by the "non-visible" engineering, data, and governance work required to build a reliable and secure AI agent for a high-stakes compliance environment. Unlike traditional software, an AI MVP's effort is front-loaded in data and model validation, not feature development.

The key drivers consuming the co-innovation budget are:

1. **Secure Data Integration & Preparation** (The "Data" Challenge):
    - Secure Pipelines: Establishing a secure, auditable, and performant data pipeline from a GCP UAT environment to the FCRM system is a major task. This involves complex networking, VPC-SC, and IAM configurations.
    - Data Quality & Transformation: FCRM data is unlikely to be perfectly clean or optimized for a GenAI model. We'll spend significant time in "build-intervene-evaluate loops" to analyze, prepare, and transform this data into a high-quality, reliable dataset that the agent can consistently use.
2. **Agent Development & Tuning** (The "AI" Challenge):
    - Prompt Engineering & Grounding: The Agent's "brain" must be meticulously built. This involves iterative prompt engineering and grounding the model in ING's specific data and business logic to ensure it generates factually accurate summaries.
3. **Reliability & Factual Accuracy**: The biggest risk in GenAI is "hallucinations" or factually incorrect outputs. A large portion of the timeline is dedicated to "validation cycles"— rigorously testing the agent's summaries, tuning the model to eliminate errors, and ensuring its accuracy is within an acceptable range for UAT.
4. **Compliance, Governance & Explainability**
    - **Audit & Explainability**: This is the most critical component. For UAT sign-off, the Agent can't be a "black box." We must build robust logging and explainability layers so that every proposal is "transparent, auditable, and consistent." An analyst (and auditor) must be able to see why the agent made its summary.
    - **Risk Mitigation**: The timeline accounts for the effort to build deterministic "guardrails" that prevent the agent from making proposals that violate defined

business rules. This validation, in partnership with ING's risk and compliance teams, is essential before any analyst can use the tool.

4   Dependencies

Successful delivery requires a collaborative partnership with clear responsibilities.

4.1.1   Dependencies on ING

- **Product Owner**: Arthur Van Bronswijk to act as the dedicated Product Owner (PO), responsible for prioritizing the backlog, providing domain decisions and validating outcomes
- **SME Access**: Access to one to two L1 FCRM analysts for three to four hours per week (during Sprints 1–3) for initial process discovery and (during Sprints 6–8) for the formal UAT execution
- **Data Access** (Critical Path): Provision of a secure one-off load to the FCRM, available within the project's initial phase. The quality and completeness of this data will directly impact the project timeline.
- **Platform Access**: A dedicated Google Cloud Platform (GCP) project and a UAT environment provisioned with the required services (see Tech Stack) and secure network access to the FCRM data source
- **Compliance & Risk Approval**: Early engagement (from Week 1) with ING's compliance, risk and legal teams. Formal sign-off from these teams is a prerequisite to begin the UAT phase (target: Week 15)

4.1.2   Dependencies on HCLTech

- Provide the dedicated squad composition (see section 8) for the project's duration
- Manage the project budget and provide transparent burndown charts
- Lead the Agentic AI design, technical development and deployment to the agreed UAT environment

5   Key Assumptions

- A stable, dataset for the required FCRM data is available or can be made available by ING within the project timeline
- The quality and structure of the FCRM data are sufficient for GenAI analysis without requiring a multi-month data-cleansing project
- ING has the necessary Google Cloud Platform (GCP) subscriptions (e.g., Gemini Enterprise) to enable the Vertex AI Gemini APIs for the project
- The required GCP services (Vertex AI, Cloud Run) and a UAT environment are approved for use
- The agreed service hours for this co-innovation project are dedicated to our (HCLTech's) squad effort. All ING-side efforts (PO, SMEs, platform engineering, compliance review) are separate
- The scope of the MVP is strictly limited to the features described in Section 1
- ING compliance and risk teams will be engaged in Sprint 0 to align on requirements, enabling a timely review process that doesn't block the start of UAT

6 Timelines

The proposed squad (Section 8) consumes ~160 hours/week. The total co-innovation budget of 3,000 service hours thus supports an ~18–19 week project.

- **Weeks 1–2** (Sprint 0): Kick-off, finalize MVP scope, define data-access contracts, set up GCP/UAT environments and initiate compliance/risk review process
- **Weeks 3–8** (Sprints 1–3): (Data & Security) Develop secure data pipeline (the Agent's "FCRM Tool"). Perform initial data analysis, preparation, and transformation.
- **Weeks 9–14** (Sprints 4–6): (Agent & UI) Build the standalone Copilot web application (UI). Develop the core agent logic, prompts, and grounding. Begin iterative testing and tuning of AI proposals.
- **Weeks 15–17** (Sprints 7–8): (Validation & UAT) Formal UAT phase. Support L1 analysts in test execution (pending compliance sign-off). Gather feedback, fix bugs and finalize agent proposal accuracy and explainability logs.
- **Week 18** (Sprint 9): Final UAT report, demo of the validated Copilot, and handover of the final UAT-ready application and documentation.
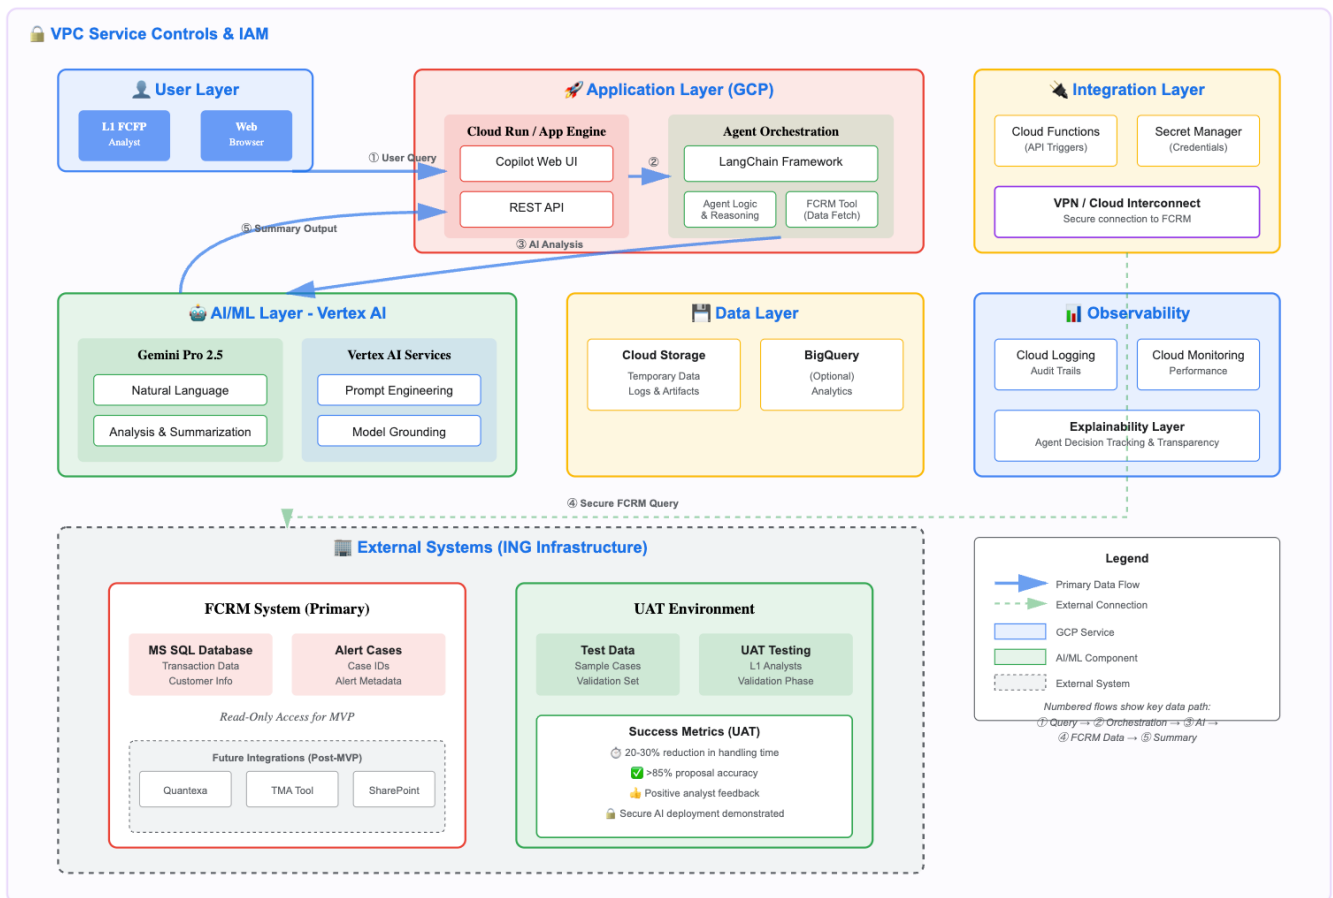
7 Success Criteria

- **Project success**
  - Deliver the defined MVP scope to the UAT environment on time and within the agreed co-innovation budget
  - Achieve formal sign-off from compliance/risk to begin the UAT phase
- **Quantitative success** (to be validated during UAT)
  - Time Savings: Achieve a 20%–30% reduction in average handling time for the in-scope alerts
  - Proposal Accuracy: The Copilot's triage proposal (Close/Escalate) is accepted by the analyst in >85% of UAT cases
- **Qualitative success** (to be validated during UAT)
  - Positive qualitative feedback from L1 analysts (e.g., "This is useful and I would use it", "Can we trust the system?")
  - Demonstrate successful, secure use of Agentic AI on ING data within the GCP UAT environment

8 Technology Stack (GCP)

- **AI/LLM**: Google Vertex AI, accessing the latest Gemini Pro models (2.5 as per October 2025)
- **Agent Framework**: A custom agent built with an open-source framework (like LangChain) to orchestrate the agent's goals, tools and reasoning. This provides the flexibility and control needed for data-residency and compliance requirements.
- **Application Hosting**: Google Cloud Run or App Engine (for the standalone web app/API)
- **Data/Integration**: Cloud Functions (for API triggers), Cloud Storage (for any temp data) and Secret Manager (for secure credentials)

- **Security & Networking**: VPC Service Controls, IAM and a secure interconnect/VPN to the ING FCRM data source
- **Database**: MS SQL

9   Squad Composition

This squad is designed for a ~19-week (4.5-month) engagement, aligned to the total co-innovation budget.

9.1   HCLTech squad (core team)

- **One** Senior AI Lead / Architect (Full-time): Leads technical design, co-manages backlog with the ING PO and performs code reviews
- **Two** AI / Machine Learning (ML) Engineers (Full-time): Build the data pipeline, AI logic and application UI/API
- **One** Scrum Master (Part-time, 50%): Facilitates ceremonies, tracks burndown of the project hours and manages dependencies

(Estimated Total Burn: ~2,660 hours over 19 weeks, leaving a ~340-hour buffer for contingency within the 3,000-hour budget.)

9.2 Required ING stakeholders

- **One** Product Owner (Arthur Van Bronswijk, Part-time): Dedicated time for backlog grooming, sprint reviews and unblocking the team – 4 hours per week. There is no need to attend daily standups.
- **Two** L1 Analyst Subject Matter Experts (SMEs) (Part-time): For domain expertise (Sprints 1–3) and formal UAT execution (Sprints 7–8)
- **One** ING Platform Engineer (As-needed): To provision the GCP/UAT environments and configure the secure data access.

HCLTech | Supercharging Progress™